

# Parallelism

CS242

Lecture 16

# Concurrency vs. Parallelism

- Concurrency
  - Multiple threads of control
  - May interleave at any granularity
    - Above atomic operations
  - Makes sense even on a single processor
  - A structuring device to deal with unpredictable latencies
- Parallelism
  - Multiple threads of control that execute at the same time
  - On multiple hardware devices
  - Inherently about performance

# Amdahl's Law


$$\text{Speedup} = \frac{1}{(1 - p) + (p / s)}$$

where

$p$  = portion of the program sped up

$s$  = factor improvement of that portion

# Parallelism: Speed vs. Processors



# Discussion

- Amdahl's law is simple and general
  - Not about a specific machine or program
- And unforgiving
  - To speed up by 1000x, must parallelize 99.9%
  - To reach 10,000x, must parallelize 99.99%
  - And these are not very aggressive targets!


# Three Topics

- Bulk synchronous model
- SPMD
- MapReduce

# Bulk Synchronous Model

- A model
  - An idealized machine
- Originally proposed for analyzing parallel algorithms
  - Leslie Valiant
  - “A Bridging Model for Parallel Computation”, 1990

# The Machine





*What are some properties of this machine model?*

# Computations

- A sequence of *supersteps*:
- Repeat:
- *All processors do local computation*
- *Barrier – all threads must reach the barrier before any proceed*
- *All processors communicate*
- *Barrier*

*What are properties of this computational model?*

# Basic Properties

- Uniform
  - compute nodes
  - communication costs
- Separate communication and computation
- Synchronization is global

# The Idea

- Programs are
  - written for  $v$  virtual processors
  - run on  $p$  physical processors
- If  $v \geq p \log p$  then
  - Managing memory, communication and synchronization can be done automatically within a constant factor of optimal

# How Does It Work?

- Roughly
  - Memory addresses are hashed to a random location in the machine
  - Guarantees that on average, memory accesses have the same cost
  - The extra  $\log p$  factor of threads are multiplexed onto the  $p$  processors to hide the latency of memory requests
  - The processors are kept busy and do no more compute than necessary

# SPMD

# Terminology

- SIMD
  - Single Instruction, Multiple Data
- SPMD
  - Single Program, Multiple Data



# SIMD = Vector Processing

```
if (factor == 0)
```

```
    factor = 1.0
```

```
A[1..N] = B[1..N] * factor;
```

```
j += factor;
```

# Picture

```
if (factor == 0)  
    factor = 1.0
```

---

```
A[1] = B[1] *  
    factor
```

```
A[2] = B[2] *  
    factor
```

```
A[3] = B[3] *  
    factor
```

```
...
```

---

```
j += factor
```

---

# Comments

- Single thread of control
  - Global synchronization at each program instruction
- Can exploit fine-grain parallelism
  - Assumption of hardware support

# SPMD = Single Program, Multiple Data

## SIMD

```
if (factor == 0)
    factor = 1.0
A[1..N] = B[1..N] * factor;
j += factor;
...
```

## SPMD

```
if (factor == 0)
    factor = 1.0
A[myid] = B[myid] * factor;
j += factor;
...
```

# Picture

```
if (factor == 0)  
    factor = 1.0
```

```
A[1] = B[1] *  
    factor
```

```
j += factor
```

```
if (factor == 0)  
    factor = 1.0
```

```
A[2] = B[2] *  
    factor
```

```
j += factor
```

...

# Comments

- Multiple threads of control
  - One (or more) per processor
- Asynchronous
  - All synchronization is programmer-specified
- Threads are distinguished by **myid**
- Choice: Are variables local or global?

# Comparison

- SIMD
  - Designed for tightly-coupled, synchronous hardware
  - i.e., vector units
- SPMD
  - Designed for clusters
  - Too expensive to synchronize every statement
  - Need a model that allows asynchrony

# MPI

- Message Passing Interface
  - A widely used standard
  - Runs on everything
- Some similarities to the Pi Calculus
  - For one-to-one communication
  - Also has one-to-all and all-to-all communication primitives
- Most popular way to write SPMD programs



# MPI Programs

- Standard sequential programs
  - All variables are local to a thread
- Augmented with calls to the MPI interface
  - SPMD model
  - Every thread has a unique identifier
  - Threads can send/receive messages
  - Synchronization primitives

# MPI Point-to-Point Routines

- `MPI_Send(buffer, count, type, dest, ...)`
- `MPI_Recv(buffer, count, type, source, ...)`

# Example

```
for (....) {  
  // p = number of chunks of 1D grid, id = process id, h[] = local chunk of the grid  
  // boundary elements of h[] are copies of neighbors boundary elements  
  
  .... Local computation ...  
  
  // exchange with neighbors on a 1-D grid  
  if ( 0 < id )  
    MPI_Send ( &h[1], 1, MPI_DOUBLE, id-1, 1, MPI_COMM_WORLD );  
  if ( id < p-1 )  
    MPI_Recv ( &h[n+1], 1, MPI_DOUBLE, id+1, 1, MPI_COMM_WORLD, &status );  
  if ( id < p-1 )  
    MPI_Send ( &h[n], 1, MPI_DOUBLE, id+1, 2, MPI_COMM_WORLD );  
  if ( 0 < id )  
    MPI_Recv ( &h[0], 1, MPI_DOUBLE, id-1, 2, MPI_COMM_WORLD, &status );  
  
  ... More local computation ...  
  
}
```

# MPI Collective Communication Routines

- MPI\_Barrier(...)
- MPI\_Bcast(...)
- MPI\_Scatter(...)
- MPI\_Gather(...)
- MPI\_Reduce(...)

# Typical Structure

```
communicate_get_work_to_do();  
barrier;           // not always needed  
do_local_work();  
barrier;  
communicate_write_results();
```

*What does this remind you of?*

# Bulk Synchronous/SPMD Model

- Easy to understand
- Phase structure guarantees no data races
  - Barrier synchronization also easy to understand
- Fits many problems well

# PGAS Model

- PGAS = *Partitioned Global Address Space*
- There is one global address space
- But each thread owns a partition of the address space that is more efficient to access
  - i.e., the local memory of a processor
- Equivalent in functionality to MPI
  - But typically presented as a programming language
  - Examples: Split-C, UPC, Titanium

# PGAS Languages

- No library calls for communication
- Instead, variables can name memory locations on other machines

// Assume *y* points to a remote location

// The following is equivalent to a send/receive

*x* = \**y*



# PGAS Languages

- Also provide collective communication
- Barrier
- Broadcast/Reduce
  - 1-many
- Exchange
  - All-to-all

# PGAS vs. MPI

- Programming model very similar
  - Both provide SPMD
- From a pragmatic point of view, MPI rules
  - Easy to add MPI to an existing sequential language
- For productivity, PGAS is better
  - Programs filled with low-level details of MPI calls
  - PGAS programs easier to modify
  - PGAS compilers can know more/do a better job

# Summary

- SPMD is well-matched to cluster programming
  - Also works well on shared memory machines
- One thread per core
  - No need for compiler to discover parallelism
  - No danger of overwhelming # of threads
- Model exposes memory architecture
  - Local vs. Global variables
  - Local computation vs. sends/receives

# MapReduce

# The Motivation

- Organizations today produce enormous amounts of data
  - Customer orders
  - Patient records
  - Social media interactions
  - Log files from all the computers
  - ...
- The only way to process this data is in parallel
  - Would take far too long on a single machine
  - Memory is the main the issue, not processing power

# The Programming Model

- Programs have a type  $\text{List}(\text{String}) \rightarrow a$
- A string here is (substantial) block of data
  - Simply interpreted as a sequence of bytes; i.e., a string
  - Example: a file consisting of tens of megabytes to terabytes
- A program has the form

$(\text{reduce } g) \diamond (\text{map } f)$

where  $f: \text{String} \rightarrow a$  and  $g: (a,a) \rightarrow a$

# Semantics

(reduce g)  $\diamond$  (map f)

What does  $\diamond$  do?

- Group the pairs returned by the map on the first components
  - the *keys*
- Apply the reduction function to the set of second components for each key
  - The *values*
- Return the set of key-value pairs after the reduction

# Example: Word Count

Count the number of occurrences of each word in a corpus.

$(\text{reduce } +) \diamond \text{map } ((\text{map } \lambda w.(w,1)) \circ \text{split } '')$

where  $\text{split}: \text{Char} \rightarrow \text{String} \rightarrow \text{List}(\text{String})$



# Discussion

- A combinator-based language!
  - Recall SKI
- The  $\diamond$  combinator is perhaps unexpected
  - More than function composition
- Makes it easy to write parallel, distributed computations
  - Without knowing anything about parallelism, different kinds of parallel hardware, synchronization, etc.
  - Just need to understand basic functional programming

# Discussion

- Why is MapReduce functional?
- Using state makes the most sense when there is a hardware-provided global state
  - i.e., a single address space
- But in distributed machines, there is no single address space
  - Or rather, it is expensive to provide one in software
- Functional programming makes sense regardless of the scale of the underlying machine

# Evolution

- MapReduce became very popular
  - Gave programmers an easy way to write simple computations that they otherwise couldn't write themselves
- Downside: Many algorithms can't be expressed in MapReduce
  - At least efficiently
- Spark is a well-known successor to MapReduce

# Summary

- Two very different parallel programming models
- Bulk synchronous/SPMD
  - Message-passing based
  - Run many copies of the same program in parallel
  - Threads are distinguished by a *myid* unique to each thread
  - Threads communicate via message passing
  - Used mostly for scientific programming
- MapReduce
  - Based on functional combinators with obvious parallel implementations
  - More expressive than simple composition because of the flexibility of the group-by operation that MapReduce adds between the map and reduce stages
  - Used for data analytics