

Taming the C Monster

Haskell FFI Techniques

Fraser Tweedale
@hackuador

May 22, 2018

Terminal

NOTMUCH(1)

notmuch

NOTMUCH(1)

NAME

notmuch - thread-based email index, search, and tagging

SYNOPSIS

notmuch [option ...] command [arg ...]

DESCRIPTION

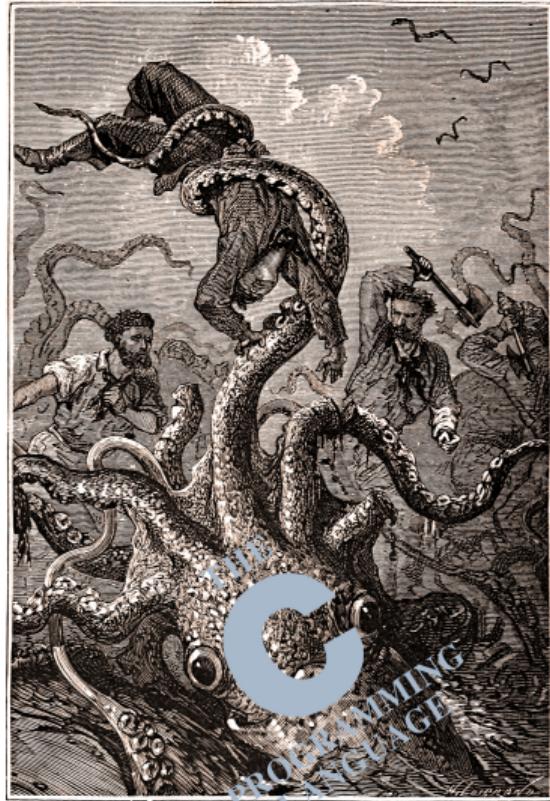
Notmuch is a command-line based program for indexing, searching, reading, and tagging large collections of email messages.

This page describes how to get started using notmuch from the command line, and gives a brief overview of the commands available. For more information on e.g. **notmuch show** consult the **notmuch-show(1)** man page, also accessible via **notmuch help show**

The quickest way to get started with Notmuch is to simply invoke the **notmuch** command with no arguments, which will interactively guide you through the process of indexing your mail.

NOTE

While the command-line program **notmuch** provides powerful functionality,
Manual page notmuch(1) line 1 (press h for help or q to quit)





File Edit View Terminal Tabs Help

Fraser Tweedale	04/Nov	Re: [purebred-mua/purebred] Only show top level messages from threads (#103)	inbox
Fraser Tweedale	04/Nov	Re: [purebred-mua/purebred] Re-factoring the event handler (#95)	inbox
Fraser Tweedale	04/Nov	Re: [purebred-mua/hs-notmuch] API to retrieve top level messages (#17)	inbox
Fraser Tweedale	31/Oct	Re: [purebred-mua/purebred] Mail add tags (#88)	archive
Simon Baird <no	30/Oct	[simonbaird/gerrit-nag] Use new reviewedby Gerrit search param (e772817)	inbox
Fraser Tweedale	27/Oct	Re: [purebred-mua/purebred] Configure "Identities" to compose an email (#93)	archive
Fraser Tweedale	25/Oct	Re: [purebred-mua/purebred] Consider listing purebred in the Brick featured projects list? (#8	
Jonathan Daughe	25/Oct	Re: [jtdaugherty/brick] Brick.Widgets.List.listReplace has with no selection when used with li	
Jonathan Daughe	25/Oct	Re: [purebred-mua/purebred] Consider listing purebred in the Brick featured projects list? (#8	
Jonathan Daughe	25/Oct	Re: [purebred-mua/purebred] Consider listing purebred in the Brick featured projects list? (#8	
Purebred:	Item 3 of 286		
from	Fraser Tweedale <notifications@github.com>		
to	purebred-mua/purebred <purebred@noreply.github.com>		
subject	Re: [purebred-mua/purebred] Mail add tags (#88)		

Merged #88.

--
You are receiving this because you authored the thread.

Reply to this email directly or view it on GitHub:

<https://github.com/purebred-mua/purebred/pull/88#event-1318097299>

FFI basics

why FFI?

- ▶ want to do \$THING in Haskell
- ▶ there exists a C library for \$THING
- ▶ interoperability / bug-compatibility
- ▶ performance / timing-critical code

C FFI

```
{-# LANGUAGE ForeignFunctionInterface #-}

import Foreign.C.Types

foreign import ccall "math.h sin"
    c_sin :: CDouble -> CDouble

main :: IO ()
main = print $ c_sin 1.0
```

`hsc2hs`

- ▶ file extension: `.hsc`
- ▶ part of GHC distribution
- ▶ good support for marshalling structs

c2hs

- ▶ file extension: .chs
- ▶ more features than hsc2hs
- ▶ automatic generation of foreign import declarations

```
library
```

```
...
```

```
build-tools:
```

```
    c2hs >= 0.19.1
```

c2hs - example

```
...
result <- {#call notmuch_database_open #} path 1 ptr
...
```

c2hs - example

```
...
    result <- notmuch_database_open path 1 ptr
    ...
foreign import ccall "Notmuch/Binding.chs.h notmuch_database_open"
    notmuch_database_open
        :: CString -> CInt -> Ptr (Ptr Database) -> IO CInt
```

Foreign.Ptr

```
data Ptr a

nullPtr :: Ptr a

plusPtr :: Ptr a -> Int -> Ptr b

castPtr :: Ptr a -> Ptr b
```

Foreign.ForeignPtr

```
data ForeignPtr a

type FinalizerPtr a = FunPtr (Ptr a -> IO ())

newForeignPtr :: FinalizerPtr a -> Ptr a -> IO (ForeignPtr a)

withForeignPtr :: ForeignPtr a -> (Ptr a -> IO b) -> IO b
```

Foreign.C.String

```
type CString = Ptr CChar
```

```
peekCString :: CString -> IO String
```

```
withCString :: String -> (CString -> IO a) -> IO a
```

Foreign.Storable

```
class Storable a where
    peek :: Ptr a -> IO a
    ...
instance Storable (Ptr a)

-- Foreign.Marshal.Alloc
alloca :: Storable a => (Ptr a -> IO b) -> IO b
```

C constructions and idioms

enum types

```
typedef enum _notmuch_status {
    NOTMUCH_STATUS_SUCCESS = 0,
    NOTMUCH_STATUS_OUT_OF_MEMORY,
    NOTMUCH_STATUS_READ_ONLY_DATABASE,
    NOTMUCH_STATUS_UNBALANCED_FREEZE_THAW,
    ...
} notmuch_status_t
```

enum types

```
{#enum notmuch_status_t as Status {underscoreToCase} deriving (Eq) #}
```

enum types

```
data Status = StatusSuccess
    | StatusOutOfMemory
    | StatusReadOnlyDatabase
    | StatusUnbalancedFreezeThaw
    ...
deriving (Eq)

instance Enum Status where
    ...
```

opaque pointer types

```
typedef struct _notmuch_database notmuch_database_t;
```

opaque pointer types

```
{#pointer *notmuch_database_t as DatabaseHandle foreign newtype #}
```

opaque pointer types

```
newtype DatabaseHandle = DatabaseHandle (ForeignPtr DatabaseHandle)

withDatabaseHandle
  :: DatabaseHandle -> (Ptr DatabaseHandle -> IO b) -> IO b
withDatabaseHandle (DatabaseHandle fptr) =
  withForeignPtr fptr
```

double-pointer constructors

```
notmuch_status_t  
notmuch_database_open (const char *path,  
                      notmuch_database_mode_t mode,  
                      notmuch_database_t **database);
```

double-pointer constructors

```
databaseOpen :: CString -> IO (Either Status DatabaseHandle)
databaseOpen path =
    alloca $ \ptr -> do
        result <- {#call notmuch_database_open #} path 1 ptr
        case toEnum (fromIntegral result) of
            StatusSuccess ->
                Right . DatabaseHandle <$> (peek ptr >>= newForeignPtr_)
            e ->
                pure (Left e)
```

iterator

```
notmuch_tags_t *
notmuch_message_get_tags (notmuch_message_t *message);

notmuch_bool_t
notmuch_tags_valid (notmuch_tags_t *tags);

const char *
notmuch_tags_get (notmuch_tags_t *tags);

void
notmuch_tags_move_to_next (notmuch_tags_t *tags);
```

iterator

```
tagsToList :: Tags -> IO [String]
tagsToList (Tags ptr) = go
  where
    go = test ptr >>= \valid -> case valid of
      0 -> pure []
      _ -> (():
              <$> (get ptr >>= mk >>= \x -> next ptr $> x)
              <*> go)

    test = {#call notmuch_tags_valid #}
    get = {#call notmuch_tags_get #}
    next = {#call notmuch_tags_move_to_next #}
    mk = peekCString
```

macros

```
void *talloc_steal(const void *new_ctx, const void *ptr);
```

macros

```
#if (_GNUC__ >= 3)
#define _TALLOC_TYPEOF(ptr) __typeof__(ptr)
#define talloc_steal(ctx, ptr) ({ \
    _TALLOC_TYPEOF(ptr) __talloc_steal_ret = (_TALLOC_TYPEOF(ptr)) \
        _talloc_steal_loc((ctx), (ptr), __location__); \
    __talloc_steal_ret; })
#else /* __GNUC__ >= 3 */
#define _TALLOC_TYPEOF(ptr) void *
#define talloc_steal(ctx, ptr) \
    (_TALLOC_TYPEOF(ptr)) _talloc_steal_loc((ctx), (ptr), __location__)
#endif /* __GNUC__ >= 3 */
void *_talloc_steal_loc(
    const void *new_ctx, const void *ptr, const char *location);
```

macros

Two options:

- ▶ bind to non-public API (e.g. `_talloc_steal_loc`)
- ▶ write “*c bits*”

external object lifecycles

```
notmuch_query_t *
notmuch_query_create (notmuch_database_t *database,
                      const char *query_string);

void
notmuch_query_destroy (notmuch_query_t *query);
```

external object lifecycles

```
query_create :: DatabaseHandle -> String -> IO (Query a)
query_create db s = withCString s $ \s' ->
    withDatabaseHandle db $ \db' ->
        {#call notmuch_query_create #} db' s'
        >>= fmap Query . newForeignPtr query_destroy

foreign import ccall "&notmuch_query_destroy"
    query_destroy :: FinalizerPtr Query
```

external object lifecycles

```
query_create :: DatabaseHandle -> String -> IO (Query a)
query_create db s = withCString s $ \s' ->
    withDatabaseHandle db $ \db' ->
        {#call notmuch_query_create #} db' s'
        >>= fmap Query . newForeignPtr query_destroy

foreign import ccall "&notmuch_query_destroy"
    query_destroy :: FunPtr (Ptr Query -> IO ())
```

external object lifecycles - beware

- ▶ hidden references in derived objects
- ▶ fancy allocators (e.g. *talloc*)

API safety

read-only mode

```
/* can return NOTMUCH_STATUS_READ_ONLY_DATABASE */
notmuch_status_t
notmuch_message_add_tag (notmuch_message_t *message, const char *tag);
```

read-only mode

```
{#enum notmuch_database_mode_t as DatabaseMode {underscoreToCase} #}
```

read-only mode

```
data DatabaseMode = DatabaseModeReadOnly  
                  | DatabaseModeReadWrite  
  
instance Enum DatabaseMode where  
  ...
```

read-only mode

```
{-# LANGUAGE DataKinds #-}

newtype Database (a :: DatabaseMode) = Database DatabaseHandle

withDatabase :: Database a -> (Ptr DatabaseHandle -> IO b) -> IO b
withDatabase (Database dbh) = withDatabaseHandle dbh

data Message (a :: DatabaseMode) = Message MessageHandle
```

read-only mode

```
type RW = 'DatabaseModeReadWrite -- convenient alias

messageAddTag :: Message RW -> Tag -> IO ()
messageAddTag msg tag = void $ withMessage msg $
    tagUseAsCString tag . {#call notmuch_message_add_tag #}
```

locking

```
/* can return NOTMUCH_STATUS_READ_ONLY_DATABASE */
notmuch_status_t
notmuch_message_freeze (notmuch_message_t *message);

/* can return NOTMUCH_STATUS_READ_ONLY_DATABASE
   or NOTMUCH_STATUS_UNBALANCED_FREEZE_THAW      */
notmuch_status_t
notmuch_message_thaw (notmuch_message_t *message);
```

locking

```
{-# LANGUAGE DataKinds #-}
{-# LANGUAGE TypeFamilies #-}
{-# LANGUAGE TypeOperators #-}

import GHC.TypeLits

data Message (n :: Nat) (a :: DatabaseMode) = Message MessageHandle

messageAddTag :: Message n RW -> Tag -> IO ()
messageAddTag msg tag = void $ withMessage msg $
    tagUseAsCString tag . {#call notmuch_message_add_tag #}
```

locking

```
messageFreeze :: Message n RW -> IO (Message (n + 1) RW)
messageFreeze msg =
    withMessage msg {#call notmuch_message_freeze #} $> coerce msg
```

```
messageThaw :: (1 <= n) => Message n RW -> IO (Message (n - 1) RW)
message_thaw msg =
    withMessage msg {#call notmuch_message_thaw #} $> coerce msg
```

Performance

unsafe

```
{#call      notmuch_messages_valid #}

foreign import ccall      "notmuch.h notmuch_messages_valid"
    notmuch_messages_valid :: Messages -> IO CInt
```

unsafe

```
{#call unsafe notmuch_messages_valid #}

foreign import ccall unsafe "notmuch.h notmuch_messages_valid"
    notmuch_messages_valid :: Messages -> IO CInt
```

unsafe

Before:

```
total time = 6.53 secs (6530 ticks @ 1000 us, 1 processor)
total alloc = 260,249,536 bytes (excludes profiling overheads)
```

After:

```
total time = 3.73 secs (3728 ticks @ 1000 us, 1 processor)
total alloc = 260,249,536 bytes (excludes profiling overheads)
```

lazy iteration

```
messagesToList :: Messages -> IO [Message n a]
messagesToList (Messages ptr) = go
  where
    go = test ptr >>= \valid -> case valid of
      0 -> pure []
      _ -> (:) <$> (get ptr >>= mk >>= \x -> next ptr $> x)
              <*> go
```


lazy iteration (search *, take 10, count tags)

Before:

```
total time = 1.79 secs (1795 ticks @ 1000 us, 1 processor)
total alloc = 59,500,568 bytes (excludes profiling overheads)
```

After:

```
total time = 0.07 secs (68 ticks @ 1000 us, 1 processor)
total alloc = 79,960 bytes (excludes profiling overheads)
```

lazy iteration (search *, count tags)

Before:

```
68,431,240 bytes maximum residency (9 sample(s))
```

```
total time = 8.37 secs (8370 ticks @ 1000 us, 1 processor)
total alloc = 218,627,008 bytes (excludes profiling overheads)
```

After:

```
40,965,384 bytes maximum residency (8 sample(s))
```

```
total time = 7.59 secs (7586 ticks @ 1000 us, 1 processor)
total alloc = 257,666,440 bytes (excludes profiling overheads)
```

things that weren't covered

- ▶ foreign export (callbacks)
- ▶ (un)marshalling C structs
- ▶ other FFIs (JVM, JavaScript, ...)
- ▶ other performance tips (avoiding unnecessary copies, interning, ...)

Questions?



Except where otherwise noted this work is licensed under
<http://creativecommons.org/licenses/by/4.0/>

<https://speakerdeck.com/frasertweedale>

@hackuador

purebred-mua/hs-notmuch